

Spatial partitioning on a masked GPU is bistable, and a scheduler that measures can ride it

Draft

Claims traced to docs/claims-and-evidence.md

Abstract

We partition a single AMD GPU between concurrent diffusion tenants using CU masks and a step-granular scheduler, and we report three things the hardware does that a static cost model cannot express.

First, the co-run penalty for two same-model tenants on disjoint 16+16 masks is *bistable*: it is either $1.297\times$ or $1.949\times$ the solo step time, the state is drawn once per process, never changes within one, and is identifiable from a single step. Over 160 processes the slow state occurs in 16.9%, [11.4%, 23.6%], and the rate is not stable between campaigns. No predictor of the draw has survived; five have been proposed and retracted.

Second, partitioning *mismatched* tenants is almost free. SDXL beside CogVideoX-2b costs 1.00–1.06 per side at every split, and both tenants advance faster than under an equal share of the whole die: at 4+28, +71.4% and +1.4%; at 28+4, +3.6% and +30.9%. The split decides how the gain divides, not whether there is one.

Third, a negative about our own design. A dual-ledger probe that charges from measurement appeared to reduce urgent SLO miss rate by 13% under same-model co-location, by cutting expensive serial fallbacks by 73–81%. Those fallbacks come from a conservative envelope we added. With the envelope removed, every policy improves and the probe’s benefit vanishes: +0.41%, [−0.0211, +0.0256]. The probe was a workaround for a mechanism we introduced.

We also report a pre-registered negative. On the mismatched workload the grid was designed around — 405 cells, nine policies — the method is 0.56% *worse* than the strongest baseline against a 20% target, because mismatched tenants do not contend and a contention manager has nothing to manage.

Finally, we describe a measurement error worth more than any of the above: a deferred-event instrument produced a two-episode transient with a quantitative account fitting five splits to within 3–9% over a $3\times$ range, and the account was fitting a stale variable. What caught it was running two instruments over the same steps in the same process, where one reported ten times the other inside an interval that physically contained it.

1 Introduction

Serving a latency-sensitive image request beside a long-running video generation on one GPU is a scheduling problem with an unusual property: the two tenants can be given *disjoint sets of compute units* rather than alternating slices of time. On AMD hardware this is `hipExtStreamCreateWithCUMask`; on NVIDIA it is the SM mask. Spatial partitioning is attractive because neither tenant has to be suspended, and a suspended diffusion step is not free to resume.

The premise of any such scheduler is a cost model: how much does a step cost at q units, and how much more does it cost when a peer holds the other $32 - q$? This paper is largely about how badly that second question behaves, and what a scheduler has to do about it.

state	processes	mean	range
fast	23 of 30	1.297	1.287–1.303
slow	7 of 30	1.949	1.923–1.962

Table 1: Co-run cost as a multiple of the solo step, one campaign. The clusters do not overlap.

Contributions.

- (1) A measurement of the co-run penalty on gfx1201 showing it is bistable and latched per process (§3).
- (2) A measurement showing mismatched tenants partition almost for free, with both sides gaining over rotation at every split (§4).
- (3) Two negatives, both about our own design: a probe whose measured benefit is entirely the avoidance of a conservative envelope we added (§5), and a pre-registered failure on the workload the grid was built around (§6).
- (4) A methodological account of an instrument that produced a beautiful and wrong quantitative model (§7).

2 Setup

One Radeon AI PRO R9700 (gfx1201, 32 maskable units, 34.2 GB). SDXL at 768×768 ; CogVideoX-2b at 480×720 , 9 frames. A step-granular executor exposes `prepare`, `run_step`, `suspend`, `resume` and `finalize`; suspension captures the latent, the step index, the scheduler’s internal cursor and the generator state. Interrupting a request between every step and re-granting it at a different mask width produces a latent identical, byte for byte, to the uninterrupted baseline’s.

Masks are installed per stream and read back every time. A stream that quietly carried the whole die would produce an unusually *low* externality, which reads as good news, so the readback is evidence rather than a sanity check.

Measuring a step. Step cost is device time between two CUDA events bracketing the step on its own stream. Reading those events with a per-step synchronize drains the pipeline and charges the drain to the measurement: it made a 0.1155 s step read as 0.256 s. Reading them one step late avoids the drain but introduces a failure discussed in §7.

3 The co-run penalty is bistable

Two SDXL tenants on disjoint 16+16 masks pay one of two costs.

Three properties make this actionable rather than merely annoying.

It latches. No process was observed in both states: 0 of 8 processes changed across nine later episodes each, although every episode built fresh adapters and re-acquired its streams. Re-forming a pairing does not redraw.

split (sdxl+cog)	SDXL ext	Cog ext	SDXL vs rot.	Cog vs rot.
4+28	1.025	1.063	+71.4%	+1.4%
8+24	1.029	1.052	+65.3%	+8.5%
16+16	1.016	1.047	+42.2%	+22.8%
24+8	1.004	1.006	+17.7%	+29.5%
28+4	1.002	1.010	+3.6%	+30.9%

Table 2: Five processes per split, three episodes, solo and co-run both by device span; 15 measurements per entry, every range within 0.005. Rotation gives a tenant the whole die for its share of the wall clock.

It is per mask pair. Interposing an 8+24 pairing never disturbed the 16+16 state (16 of 16 checks), and 8+24 carried its own independent draw.

One step identifies it. Comparing each episode’s first step against its own median gives a mean absolute error of 0.00% over 72 episodes; the per-step series inside an episode is constant.

Pooled over 160 processes the slow state occurs in 27, 16.9%, Clopper–Pearson [11.4%, 23.6%]. The rate is *not* stable across campaigns — 22.7%, 35.0%, 8.3%, 11.7%, with Fisher $p = 0.025$ between the extremes — and we report no predictor. Five were proposed and retracted: working set, power cap, uncontrollable bistability, launch stagger, and prior use of a full-die mask.

It is RDNA4’s, not AMD’s. On an MI250X (gfx90a, CDNA2, 104 units) the same measurement — same instrument, same model, the same half-and-half split at 52+52, 32 processes rotating across four GCDs — gives 1.2336 with a standard deviation of 0.0030 and nothing above the slow threshold. At the gfx1201 rate that has probability 0.0027. The four GCDs do differ from one another, deterministically, by 0.6% with non-overlapping ranges, so the instrument resolves far less than a 50% bimodality on that machine. Partitioning still pays there, +5.6% against rotation. The mechanism generalises; the pathology does not.

The consequence for a cost model is direct. Against rotation at 2×108.9 ms, the fast state makes partitioning worth +11.5% and the slow state worth –25.5%. A model with one externality per width pair is calibrating against a draw it cannot see.

4 Mismatched tenants partition almost for free

Both tenants gain at every split, so partitioning Pareto-dominates rotation here. Neither is paying for the other’s gain — which the same-model case never manages — and no draw is observed: the two arrangements behave differently and the die keys its state on the mask pair.

5 A scheduler that measures, and what it was really doing

The runtime charges each tenant from the *measured* cost of the step it just ran, never from the cost model. Charging from prediction would make the accounting exact by construction and blind. The policy pairs tenants by matched step cost and probes: if a formed

policy	envelope on	envelope off
static SM partition	0.3011	0.2415
EDF	0.2999	0.2574
step-matched pairing	0.3011	0.2415
probe	0.2619	0.2353

Table 3: Urgent SLO miss rate, 9 paired seeds. Every policy improves with the envelope off; all four intervals exclude zero.

pairing’s observed step exceeds its paired expectation by a tolerance, the pairing is dropped for a round. The runtime adds a safety envelope that falls back to a serial action when a pairing has no measured profile or when drift exceeds 15%.

Decomposed with three arms from one code path — pairing only, plus the deadline actions, plus the probe — paired by seed, the probe carried –13.14% in the fast state and –14.20% in the slow one, with the deadline actions null in both, and it cut serial fallbacks by 73–81%.

Then we removed the envelope. Every policy improved. And the probe’s advantage over step-matched pairing went from –13.01%, [-0.0668, -0.0132], to +0.41%, [-0.0211, +0.0256] — an interval spanning zero with the point estimate on the wrong side.

So the mechanism was identified correctly and its significance was backwards. The probe does cut serial fallbacks; those fallbacks are harmful; and they are harmful because they are ours. The envelope fired on drift measured as $|\text{observed}/\text{expected} - 1|$, which treats a pairing arriving *cheaper* than predicted as a reason to take the whole die serial. Correcting it to fire on over-run only removed half the firings, and removing it entirely was better still. The envelope’s other branch, a fallback when the cost model has no entry for a pairing, never fired in any workload we ran, so it is untested rather than refuted.

The best configuration we measured is the envelope off and no probe: step-matched pairing at 0.2276 against 0.2286 with the probe.

Robustness. Injecting predictor error into the belief the policy reads, and nowhere else, gives zero safety failures at $\pm 20\%$ across 134 cells with the envelope both on and off, where a safety failure is defined in code as over-committing the die, charging from the cost model when a measurement existed, or charging a measurement taken at another quota.

6 A pre-registered negative

On the mismatched workload — 405 cells, nine policies, five seeds, loads {0.6, 0.85, 1.05} and bursts {2, 4, 8} — the method is 0.56% *worse* than the strongest non-oracle baseline, [+0.0000, +0.0056] over 45 paired configurations, against a target of a 20% improvement.

It follows from §4. Mismatched pairings cost 1.02–1.06, so there is nothing for a contention manager to manage, while a spurious probe still costs a good pairing. The pairing itself does win: FCFS by 6.5%, static partitioning and EDF by about 3%, and at overload FCFS by 17%.

How much of the grid could separate anything. “45 paired configurations” overstates the evidence, so we say what it is. The probe and step-matched pairing produce *identical* miss rates in 42 of the 45; three differ, all in the same direction. A bootstrap resample that draws none of the three has probability $(42/45)^{45} = 4.5\%$, above 2.5%, so the interval’s lower bound is zero as arithmetic and cannot be read as significance — the grid shows the probe does not help, not that it hurts. The comparisons above are wider: pairing differs from FCFS in 21 configurations, from static partitioning in 16, from EDF in 17. Two tenants are both runnable for only 26.1% of the horizon at load 0.6 and 44.9% at 1.05, and three of the 45 configurations draw no video tenant at all, so most of this grid could not have separated any two partitioning policies whatever they did.

We state this as a negative because it was registered as one. The comparison target was fixed by rule — best baseline mean, computed not chosen — because the same method beats FCFS by 30%, the strongest baseline by 22%, and its own ablation by 4.5% on a single cell, and which opponent a headline uses is the easiest thing to lose between a table and a sentence.

7 The instrument that fitted an artifact

Reading a step’s events one step late avoids draining the pipeline, but the read only lands once the previous step’s events have retired. When the measured side’s steps are short it never lands, and the variable holds one stale value for an entire episode.

This produced a two-episode “serialisation transient” in which each side appeared to cost the sum of the two solo steps. The account fitted five splits to within 3–9% over a 3× range of that sum. It was a real fit to a stale variable.

What caught it was not reasoning. Running a second instrument — events bracketing the same step on the same stream, which physically *contain* the first — over the same steps in the same process gave 1.022 where the first gave 10.888. Containment made the contradiction unarguable.

We record the working rules this produced, because the same class of error — a property of the measurement presenting itself as a property of the thing measured — occurred nine times: measure the null before varying parameters; state the discriminating criterion and its direction before running; a single co-run measurement is not evidence; the measurement architecture must match the claimed deployment architecture; and make the nuisance variable identical across arms rather than trusting it to balance — seed with position, deadline with arm, and co-run state with arm each confounded a campaign before that was routine.

8 Threats to validity

Single SKU. Every result is from one mask implementation on one card. Nothing here rules out that they depend on that implementation’s queueing and arbitration, and the bistability of §3 makes that a live question rather than a formality. No amount of additional cells on the same card substitutes for cross-vendor agreement.

Same-model co-location is not the primary workload. The probe results of §5 are measured there because the mechanism is inert on the mismatched one.

Small n on the slow state. Ten processes. The interval excludes zero; it is not tight.

An unexplained ablation. Making the runtime blind to the externality term *reduced* serial fallbacks where the mechanism and a unit test both say it should increase them. Under investigation.

9 Conclusion

The claim this work set out to make — that spatial partitioning raises utilisation by a fixed factor — is false as stated, and the reason is interesting: the factor is not fixed. It is bistable, latched per process, and depends on whether the co-located tenants contend at all.

What survives is a hardware result and a method for getting it. What does not is our own scheduling contribution: the probe’s measured value was the avoidance of an envelope we had added, and the envelope was a net cost. We report both because a scheduler that looks good against a runtime carrying one of your own defects looks good for the wrong reason, and the only way to find that out is to remove the defect and measure again.